JuliaGPs + Turing.jl

Will Tebbutt

2023-09-22

Who am I?

- Currently a postdoc with Hong
- I work on AD in Julia
- I used to work on GPs (approximate inference and software)

Outline

- GP refresher
- A complete example using JuliaGPs + Turing.jl
- A dive into the design of JuliaGPs
  - KernelFunctions.jl
  - AbstractGPs.jl
- Further examples
- Everything will be interactive

Refresher

Refresher

$$f \sim \mathcal{GP}(0, \kappa_\theta)$$

Refresher

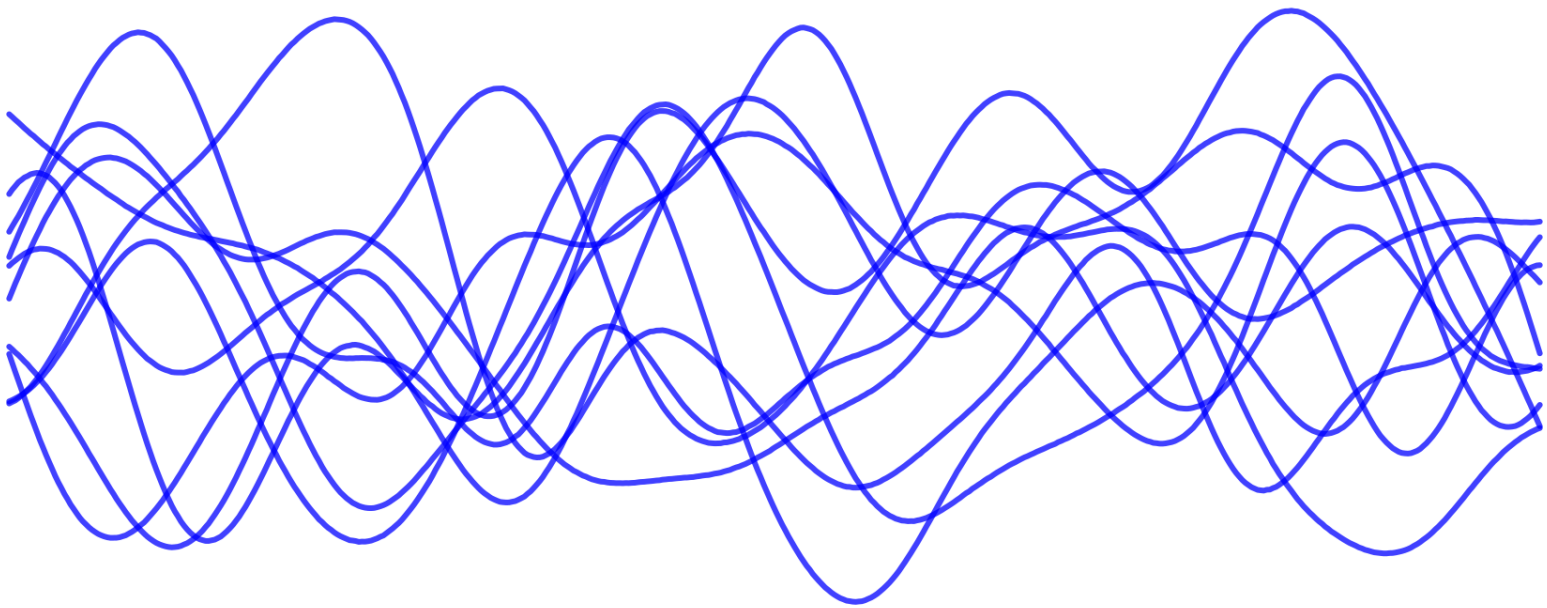

Refresher

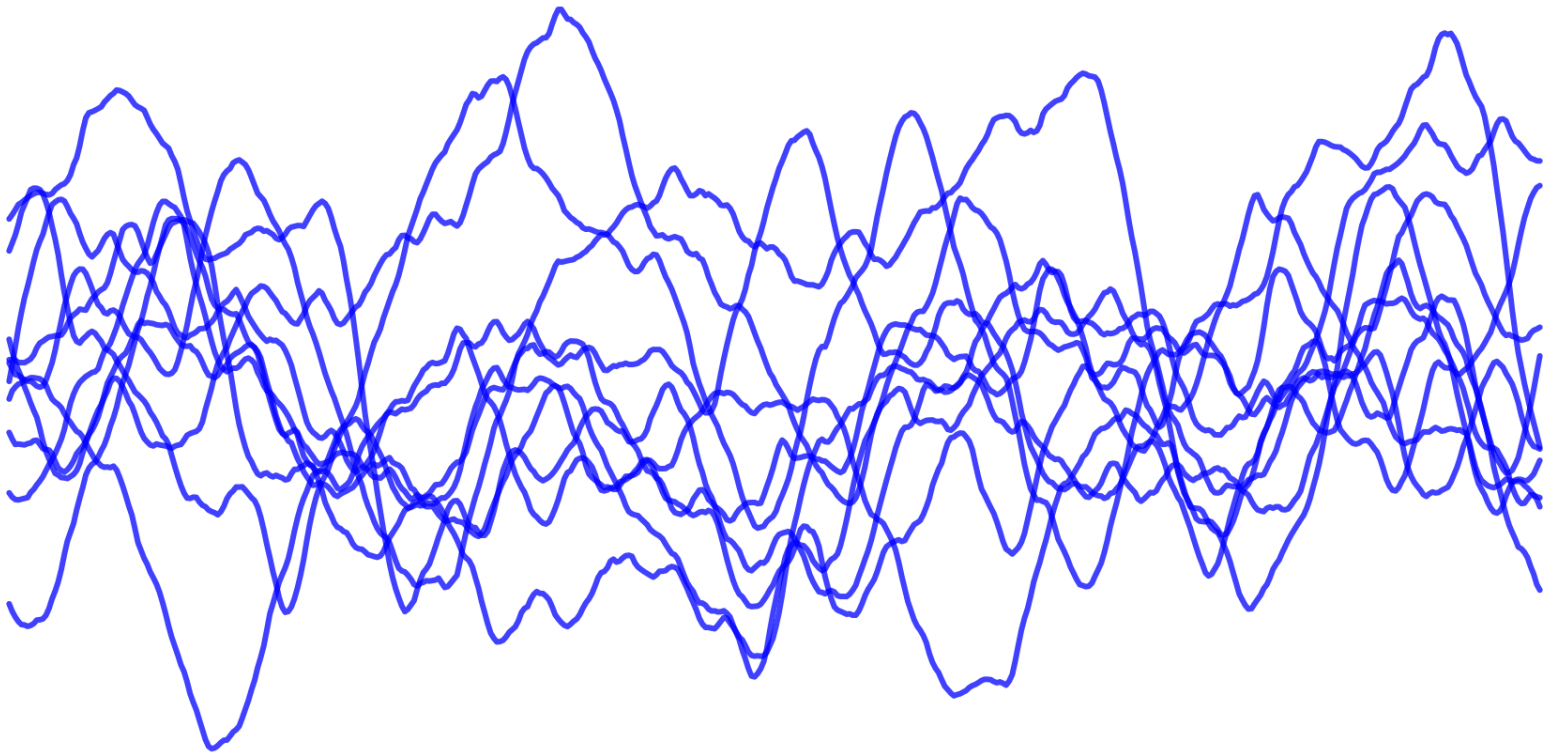

Refresher

$$f \sim \mathcal{GP}(0, \kappa_\theta)$$

Refresher

$$f \sim \mathcal{GP}(0, \kappa_\theta)$$

$$\mathbf{f} := [f(x_1), ..., f(x_J)]^\top$$

Refresher

$$f \sim \mathcal{GP}(0, \kappa_\theta)$$

$$\mathbf{f} := [f(x_1), ..., f(x_J)]^\top$$

$$\mathbf{y} \mid \mathbf{f} \sim \mathcal{N}\left(\mathbf{f}, \sigma^2 \mathbf{I}\right)$$

Refresher

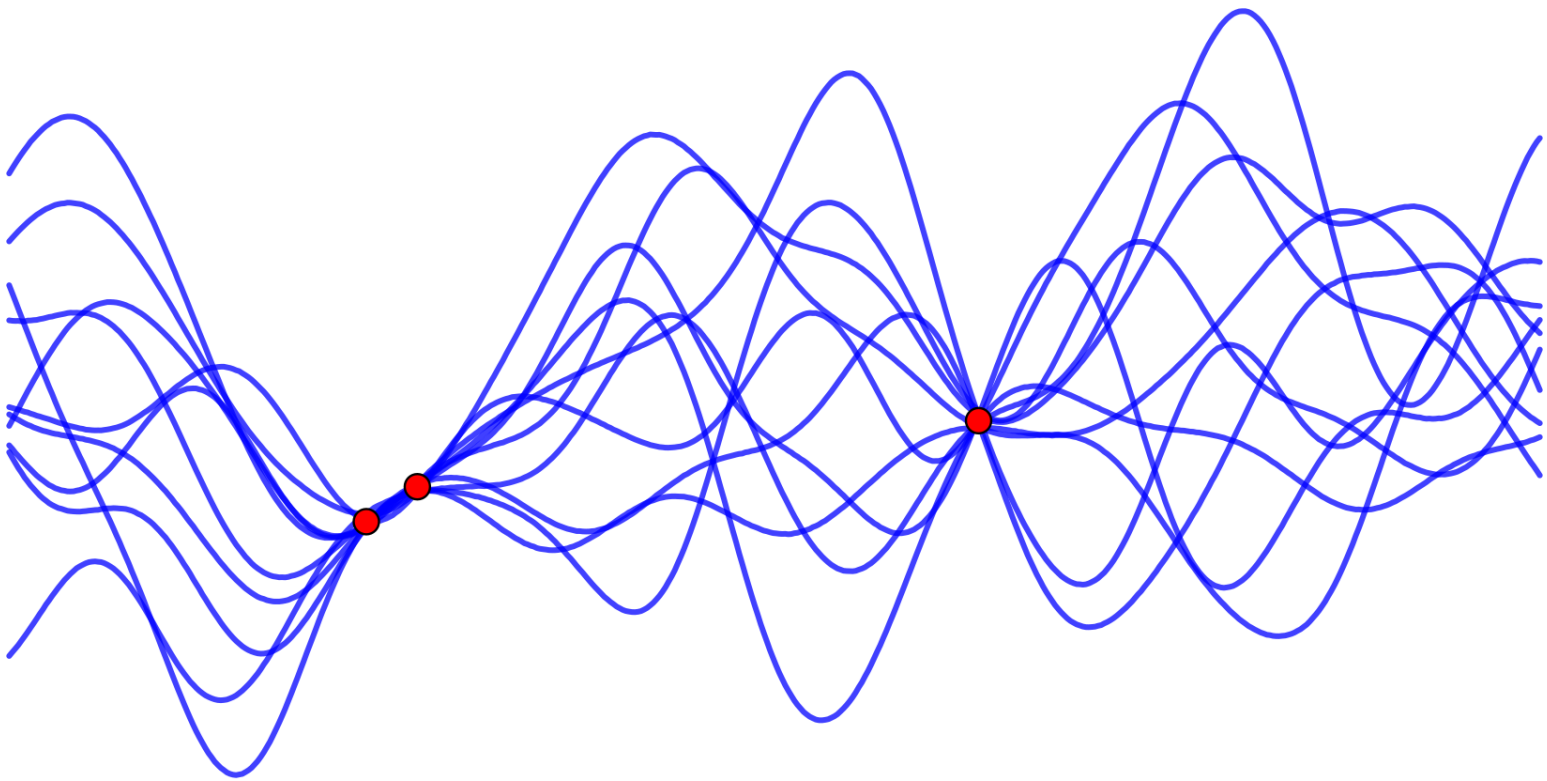

Refresher

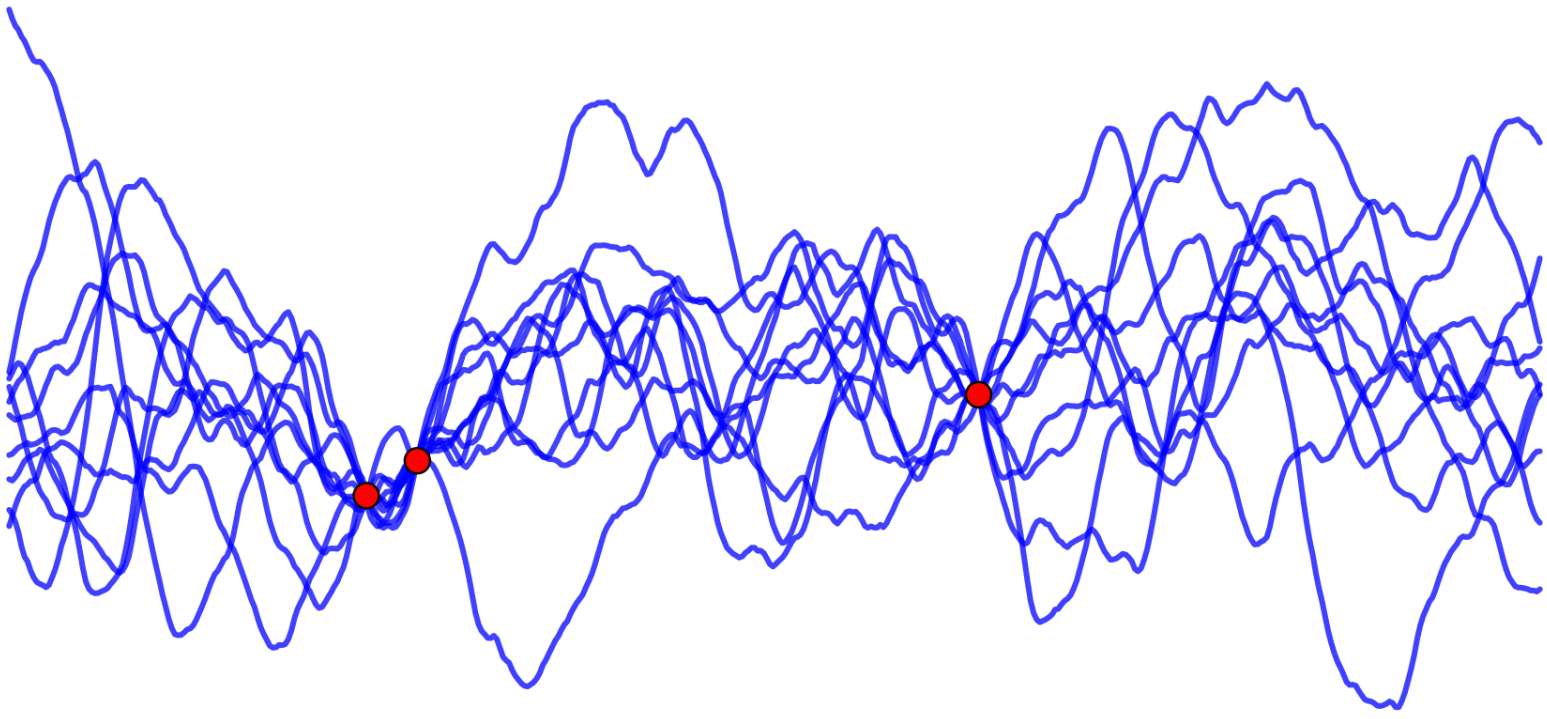

Running Toy Example

- Putting example from BDA (Gelman et al, 1995)
- *Incredibly* simple
- Non-Gaussian
- Small data

Running Toy Example

Row	distance Int64	n Int64	y Int64
1	2	1443	1346
2	3	694	577
3	4	455	337
4	5	353	208
5	6	272	149

Running Toy Example

$$f \sim \mathcal{GP}(0, \kappa_\theta)$$

$$\mathbf{f} := [f(x_1), ..., f(x_J)]^\top$$

Running Toy Example

$$f \sim \mathcal{GP}(0, \kappa_\theta)$$

$$\mathbf{f} := [f(x_1), ..., f(x_J)]^\top$$

$$\mathbf{y}_j \mid \mathbf{f}_j \sim \text{Binomial}(n_j, g(\mathbf{f}_j))$$

Running Toy Example

$$f \sim \mathcal{GP}(0, \kappa_\theta)$$

$$\mathbf{f} := [f(x_1), ..., f(x_J)]^\top$$

$$\mathbf{y}_j \mid \mathbf{f}_j \sim \text{Binomial}(n_j, g(\mathbf{f}_j))$$

$$g(x) := (1 + e^{-x})^{-1}$$

Running Toy Example

Prior Predictive Checks